

```
In [83]: #importing the library
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [84]: df = pd.read_csv('airquality.csv')
df
```

Out[84]:

	Unnamed: 0	Ozone	Solar.R	Wind	Temp	Month	Day
0	1	41.0	190.0	7.4	67	5	1
1	2	36.0	118.0	8.0	72	5	2
2	3	12.0	149.0	12.6	74	5	3
3	4	18.0	313.0	11.5	62	5	4
4	5	NaN	NaN	14.3	56	5	5
...	...	...	...	...	...	...	...
148	149	30.0	193.0	6.9	70	9	26
149	150	NaN	145.0	13.2	77	9	27
150	151	14.0	191.0	14.3	75	9	28
151	152	18.0	131.0	8.0	76	9	29
152	153	20.0	223.0	11.5	68	9	30

153 rows × 7 columns

```
In [85]: #dropping the unnamed column
df= df.drop(labels = ['Unnamed: 0'],axis=1)
df
```

Out[85]:

	Ozone	Solar.R	Wind	Temp	Month	Day
<b>0</b>	41.0	190.0	7.4	67	5	1
<b>1</b>	36.0	118.0	8.0	72	5	2
<b>2</b>	12.0	149.0	12.6	74	5	3
<b>3</b>	18.0	313.0	11.5	62	5	4
<b>4</b>	NaN	NaN	14.3	56	5	5
...	...	...	...	...	...	...
<b>148</b>	30.0	193.0	6.9	70	9	26
<b>149</b>	NaN	145.0	13.2	77	9	27
<b>150</b>	14.0	191.0	14.3	75	9	28
<b>151</b>	18.0	131.0	8.0	76	9	29
<b>152</b>	20.0	223.0	11.5	68	9	30

153 rows × 6 columns

```
In [86]: #calculating the null values
df.isnull().sum()
```

```
Out[86]: Ozone      37
Solar.R      7
Wind         0
Temp         0
Month         0
Day          0
dtype: int64
```

```
In [87]: column_means= df.mean()  
df = df.fillna(column_means)  
df
```


Out[87]:

	Ozone	Solar.R	Wind	Temp	Month	Day
0	41.00000	190.000000	7.4	67	5	1
1	36.00000	118.000000	8.0	72	5	2
2	12.00000	149.000000	12.6	74	5	3
3	18.00000	313.000000	11.5	62	5	4
4	42.12931	185.931507	14.3	56	5	5
...	...	...	...	...	...	...
148	30.00000	193.000000	6.9	70	9	26
149	42.12931	145.000000	13.2	77	9	27
150	14.00000	191.000000	14.3	75	9	28
151	18.00000	131.000000	8.0	76	9	29
152	20.00000	223.000000	11.5	68	9	30

153 rows × 6 columns

```
In [88]: plt.xlabel('Temp')  
plt.ylabel('Ozone')  
plt.scatter(df['Temp'],df['Ozone'])
```

Out[88]: <matplotlib.collections.PathCollection at 0x155e890e500>



```
In [89]: # data set for two variables.
df2 = df[['Ozone', 'Temp']]
df2
```

Out[89]:

	Ozone	Temp
0	41.00000	67
1	36.00000	72
2	12.00000	74
3	18.00000	62
4	42.12931	56
...	...	...
148	30.00000	70
149	42.12931	77
150	14.00000	75
151	18.00000	76
152	20.00000	68

153 rows × 2 columns

```
In [90]: df2.corr()
```

Out[90]:

	Ozone	Temp
Ozone	1.000000	0.608742
Temp	0.608742	1.000000

```
In [91]: #splitting the dataset into 75 and 20 percent
from sklearn.model_selection import train_test_split
x = df2['Temp']
Y = df2['Ozone']
X = x.array.reshape(-1,1)

X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.25)
```

```
In [92]: from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
```

```
In [93]: regressor = regressor.fit(X_train,Y_train)
regressor
```

Out[93]:

```
LinearRegression
LinearRegression()
```


```
In [94]: Y_pred = regressor.predict(X_test)
         results = pd.DataFrame({'Actual': Y_test, 'Predicted': Y_pred})
         print(results)
```

	Actual	Predicted
37	29.00000	50.568770
145	36.00000	48.746363
41	42.12931	70.615251
68	97.00000	68.792844
122	85.00000	72.437659
137	13.00000	30.522290
135	28.00000	41.456734
23	32.00000	12.298216
99	89.00000	65.148029
103	44.00000	57.858400
130	23.00000	43.279141
115	45.00000	45.101548
121	84.00000	76.082473
31	42.12931	43.279141
0	41.00000	23.232660
64	42.12931	54.213585
101	42.12931	68.792844
75	7.00000	46.923956
16	34.00000	21.410253
116	168.00000	48.746363
77	35.00000	50.568770
60	42.12931	52.391178
42	42.12931	68.792844
106	42.12931	45.101548
49	12.00000	34.167104
93	9.00000	48.746363
13	14.00000	25.055068
17	6.00000	5.008587
94	16.00000	50.568770
1	36.00000	32.344697
39	71.00000	65.148029
144	23.00000	30.522290
150	14.00000	37.811919
148	30.00000	28.699882
120	118.00000	72.437659
113	9.00000	32.344697
22	4.00000	12.298216
28	45.00000	48.746363
143	13.00000	17.765438

```
In [95]: from sklearn.metrics import r2_score
         r2_score(Y_test, Y_pred)
```

Out[95]: 0.3266134652335819

```
In [96]: plt.xlabel('Temp')
plt.ylabel('Ozone')
plt.scatter(X_train, Y_train,color='g')
plt.plot(X_test, Y_pred,color='k')
plt.show()
```



```
In [97]: from sklearn.metrics import mean_absolute_error, mean_squared_error
mae = mean_absolute_error(Y_test, Y_pred)
mse = mean_squared_error(Y_test, Y_pred)
rmse = np.sqrt(mse)
```

```
In [98]: print(f'Mean absolute error: {mae:.2f}')
print(f'Mean squared error: {mse:.2f}')
print(f'Root mean squared error: {rmse:.2f}')
```

```
Mean absolute error: 19.01
Mean squared error: 756.61
Root mean squared error: 27.51
```

```
In [99]: # Correlation between all features
df.corr()
```

Out[99]:

	Ozone	Solar.R	Wind	Temp	Month	Day
Ozone	1.000000	0.302970	-0.530936	0.608742	0.149081	-0.011355
Solar.R	0.302970	1.000000	-0.055245	0.262569	-0.072904	-0.145621
Wind	-0.530936	-0.055245	1.000000	-0.457988	-0.178293	0.027181
Temp	0.608742	0.262569	-0.457988	1.000000	0.420947	-0.130593
Month	0.149081	-0.072904	-0.178293	0.420947	1.000000	-0.007962
Day	-0.011355	-0.145621	0.027181	-0.130593	-0.007962	1.000000

```
In [100]: from sklearn.model_selection import train_test_split
X1 = df.drop(columns = 'Ozone').values
Y1 = df['Ozone'].values
X1_train,X1_test,Y1_train,Y1_test = train_test_split(X1,Y1,test_size=0.75)
```

```
In [101]: print(X1)
```

```
[ 8.      9.7      59.      5.      21.      ]
[320.      16.6      73.      5.      22.      ]
[ 25.      9.7      61.      5.      23.      ]
[ 92.      12.      61.      5.      24.      ]
[ 66.      16.6      57.      5.      25.      ]
[266.      14.9      58.      5.      26.      ]
[185.93150685  8.      57.      5.      27.      ]
[ 13.      12.      67.      5.      28.      ]
[252.      14.9      81.      5.      29.      ]
[223.      5.7      79.      5.      30.      ]
[279.      7.4      76.      5.      31.      ]
[286.      8.6      78.      6.      1.      ]

[287.      9.7      74.      6.      2.      ]
[242.      16.1      67.      6.      3.      ]
[186.      9.2      84.      6.      4.      ]
[220.      8.6      85.      6.      5.      ]
[264.      14.3      79.      6.      6.      ]
[127.      9.7      82.      6.      7.      ]
[273.      6.9      87.      6.      8.      ]
[291.      13.8      90.      6.      9.      ]
```



In [102]: `print(Y1)`

```
[ 41.      36.      12.      18.      42.12931034
 28.      23.      19.       8.      42.12931034
   7.      16.      11.     14.      18.
 14.      34.       6.     30.      11.
   1.      11.       4.     32.      42.12931034
42.12931034 42.12931034 23.      45.      115.
 37.      42.12931034 42.12931034 42.12931034 42.12931034
42.12931034 42.12931034 29.      42.12931034 71.
 39.      42.12931034 42.12931034 23.      42.12931034
42.12931034 21.      37.      20.      12.
 13.      42.12931034 42.12931034 42.12931034 42.12931034
42.12931034 42.12931034 42.12931034 42.12931034 42.12931034
42.12931034 135.      49.      32.      42.12931034
 64.      40.      77.      97.      97.
 85.      42.12931034 10.      27.      42.12931034
   7.      48.      35.      61.      79.
 63.      16.      42.12931034 42.12931034 80.
108.      20.      52.      82.      50.
 64.      59.      39.       9.      16.
 78.      35.      66.     122.      89.
110.      42.12931034 42.12931034 44.      28.
 65.      42.12931034 22.      59.      23.
 31.      44.      21.       9.      42.12931034
 45.     168.      73.     42.12931034 76.
118.      84.      85.      96.      78.
 73.      91.      47.      32.      20.
 23.      21.      24.      44.      21.
 28.       9.      13.      46.      18.
 13.      24.      16.      13.      23.
 36.       7.      14.      30.      42.12931034
 14.      18.      20.      ]
```

In [103]: `from sklearn.linear_model import LinearRegression`  
`regressor1 = LinearRegression()`  
`regressor1.fit(X1_train,Y1_train)`

Out[103]:

▼ LinearRegression

LinearRegression()

```
In [104]: Y1_pred = regressor1.predict(X1_test)
          results = pd.DataFrame({'Actual': Y1_test, 'Predicted': Y1_pred})
          print(results)
```

	Actual	Predicted
0	42.12931	52.235006
1	91.00000	67.732838
2	42.12931	61.680567
3	42.12931	34.555904
4	14.00000	44.143845
..	...	...
110	118.00000	66.189771
111	42.12931	13.439943
112	42.12931	49.717369
113	7.00000	52.265827
114	13.00000	26.079276

[115 rows x 2 columns]

```
In [105]: r2_score(Y1_test, Y1_pred)
```

```
Out[105]: 0.3815797431103587
```

```
In [106]: from sklearn.metrics import mean_absolute_error, mean_squared_error
          mae = mean_absolute_error(Y1_test, Y1_pred)
          mse = mean_squared_error(Y1_test, Y1_pred)
          rmse = np.sqrt(mse)
```

```
In [107]: print(f'Mean absolute error: {mae:.2f}')
          print(f'Mean squared error: {mse:.2f}')
          print(f'Root mean squared error: {rmse:.2f}')
```

Mean absolute error: 17.71  
Mean squared error: 527.82  
Root mean squared error: 22.97

```
In [108]: regressor1.score(X1_train, Y1_train)
```

```
Out[108]: 0.4079207012555187
```

```
In [109]: #New dataframe
df_new = df[['Ozone', 'Solar.R', 'Wind', 'Temp']]
df_new.head()
```

```
Out[109]:
```

	Ozone	Solar.R	Wind	Temp
0	41.00000	190.000000	7.4	67
1	36.00000	118.000000	8.0	72
2	12.00000	149.000000	12.6	74
3	18.00000	313.000000	11.5	62
4	42.12931	185.931507	14.3	56

```
In [110]: from sklearn.model_selection import train_test_split
X2 = df_new.drop(columns = 'Ozone')
Y2 = df_new['Ozone']
X2_train,X2_test,Y2_train,Y2_test = train_test_split(X2,Y2,test_size=0.75)
```

```
In [111]: from sklearn.linear_model import LinearRegression
reg= LinearRegression()
reg.fit(X2_train,Y2_train)
```

```
Out[111]:
```

LinearRegression

LinearRegression()

```
In [112]: reg.score(X2_train,Y2_train)
```

```
Out[112]: 0.5582358933693639
```

```
In [113]: print(reg.intercept_)
-93.11743667777014
```

```
In [114]: print(reg.coef_)
[ 0.03530763 -0.62420077  1.72068999]
```

```
In [115]: print(reg.predict([[229.000000 , 10.3 , 90]]))
[63.40084154]
```

```
C:\Users\rushi\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\
base.py:450: UserWarning: X does not have valid feature names, but LinearRegr
ession was fitted with feature names
warnings.warn(
```

```
In [116]: Y2_pred = reg.predict(X2_test)
          results = pd.DataFrame({'Actual': Y2_test, 'Predicted': Y2_pred})
          print(results)
```

	Actual	Predicted
142	16.00000	50.082370
98	122.00000	66.530615
88	82.00000	61.204722
66	40.00000	53.982639
137	13.00000	25.827698
..	...	...
67	77.00000	64.864764
139	18.00000	21.463731
65	64.00000	53.007344
42	42.12931	68.270303
36	42.12931	43.212216

[115 rows x 2 columns]

```
In [117]: r2_score(Y2_test, Y2_pred)
```

Out[117]: 0.39433768379241485

```
In [118]: from sklearn.metrics import mean_absolute_error, mean_squared_error
          mae = mean_absolute_error(Y2_test, Y2_pred)
          mse = mean_squared_error(Y2_test, Y2_pred)
          rmse = np.sqrt(mse)
```

```
In [119]: print(f'Mean absolute error: {mae:.2f}')
          print(f'Mean squared error: {mse:.2f}')
          print(f'Root mean squared error: {rmse:.2f}')
```

Mean absolute error: 16.62  
Mean squared error: 542.41  
Root mean squared error: 23.29

In [ ]:

In [ ]: